

Sparse Attention Mechanisms for Representation Learning

Abstract

This paper investigates advanced representations and algorithmic trade-offs.

1. Introduction

Modern visual systems require efficient representations that preserve fine-grained features.




Fig. 1: Attention architecture and perplexity scaling across sequence lengths.

2. Methodology and Experiments

2. Methodology

Detailed experimental evaluations across benchmarks.

3. Experiments

Detailed experimental evaluations across benchmarks.

4. Ablations

Detailed experimental evaluations across benchmarks.

Table 1: Accuracy comparisons across dense and sparse variants.

--